

COMS W 4111-002

W4111 - Introduction to Databases

Section 003/V03, Fall 2022

Take Home Final

Exam Instructions

- We will publish instructions on Ed.

Environment Setup and Test

MySQL

- Replace `root` and `dbuserdbuser` for the correct values for your MySQL instance from previous homework assignments and exams.
- You will need the [sample database](#) that comes with the recommended textbook to execute the setup test.
 - You should have already installed the database because you need for previous assignments.
 - I named my database

In [1]: `%load_ext sql`

In [2]: `%sql mysql+pymysql://root:dbuserdbuser@localhost`

In [3]: `%sql select * from COMS4111.student`

* mysql+pymysql://root:***@localhost
13 rows affected.

Out[3]:

| ID | name | dept_name | tot_cred |
|-------|----------|------------|----------|
| 00128 | Zhang | Comp. Sci. | 102 |
| 12345 | Shankar | Comp. Sci. | 32 |
| 19991 | Brandt | History | 80 |
| 23121 | Chavez | Finance | 110 |
| 44553 | Peltier | Physics | 56 |
| 45678 | Levy | Physics | 46 |
| 54321 | Williams | Comp. Sci. | 54 |
| 55739 | Sanchez | Music | 38 |
| 70557 | Snow | Physics | 0 |
| 76543 | Brown | Comp. Sci. | 58 |
| 76653 | Aoi | Elec. Eng. | 60 |
| 98765 | Bourikas | Elec. Eng. | 98 |
| 98988 | Tanaka | Biology | 120 |

Neo4j

- Please set the values for your Neo4j database below.

- Make sure that your database is active. If you have not used it for a while, you need to log in through the website and restart the database.

```
In [4]: neo4j_url = "neo4j+s://d4736a05.databases.neo4j.io"
neo4j_user = "neo4j"
neo4j_password = "KCvzV_oo2gWXDp1AfdRKBqThznmEwY5McMesMrQ6y3c"
```

```
In [5]: from py2neo import Graph
```

```
In [6]: def t1():
graph = Graph(neo4j_url, auth=(neo4j_user, neo4j_password))
q = "match (r:Person) where r.name='Tom Hanks' return r"
res = graph.run(q)

for r in res:
    print(r)
```

- Please rerun the following cell.

```
In [7]: t1()

Node('Person', born=1956, name='Tom Hanks')
```

MongoDB

- Please set your URL for MongoDB Atlas and make sure that your cluster is not suspended.

```
In [8]: mongodb_url = "mongodb+srv://lily:dbuserdbuser@cluster0.q3j5jvq.mongodb.net/?retryWrites=true&w=majority"
```

```
In [9]: import pymongo
```

```
In [10]: def connect():
client = pymongo.MongoClient(
    mongodb_url
)
return client

def t_connect():
c = connect()
print("Databases = ", list(c.list_databases()))
```

```
In [11]: #
# Note, you list of local databases will be different. The values do not matter.
#
t_connect()
```

```
Databases = [{'name': 'hw4', 'sizeOnDisk': 405504, 'empty': False}, {'name': 'testdb', 'sizeOnDisk': 73728, 'empty': False}, {'name': 'admin', 'sizeOnDisk': 344064, 'empty': False}, {'name': 'local', 'sizeOnDisk': 6464942080, 'empty': False}]
```

Written Questions — General Knowledge

- The written questions require a short, succinct answer.
- Remember, "If you can't explain it simply, you don't understand it well enough."
- Some questions will research using the web, lecture slides, etc. You cannot cut and paste from sources. Your answer must show that you read the material and understand the concept.

- If you use a source other than lecture material, please provide a URL to the source(s) you read.

G1

Question: List at least two reasons why database systems support data manipulation using a declarative query language such as SQL, instead of just providing a library of C or C++ functions to carry out data manipulation.

Answer:

1. Declarative query language lets us define what data we want to compute and lets the engine underneath take care of seamlessly retrieving it
2. Declarative language is typically easy to learn and understand, such as SELECT, INSERT INTO, UPDATE... the syntactical rules are not complex in SQL, which makes it a user-friendly language
3. Using a declarative query language may provide better code than can be produced manually

G2

Question: List four significant differences between:

- Processing data by writing programs that manipulate files.
- Using a database management system and query language.

Answer:

1. The file system has a higher rate of data inconsistency, while a database management system has relative low rate of data inconsistency
2. Redundant data is present in the file system, while in contrast there is no presence of redundant data in database management system
3. File systems do not handle complex transactions, whereas SQL makes it simple to build complex transactions in DBMS systems.
4. The file system has less security, but database management system supports more security mechanisms

G3

Question: List five responsibilities (functionality provided) of a database-management system. For each responsibility, explain the potential problems that would occur without the functionality.

Answer:

1. Concurrency: Even with effective integrity enforcement in each transaction, consistency restrictions may still be broken
2. Security: Unauthorized users may access the portion of database which they do not have permission
3. Backup and recovery: Data could be lost permanently and not able to recover to the state that existed prior to a failure
4. Integrity: constraints on consistency might not be met, e.g. account balances might fall below the minimum permitted
5. Interaction with the File Manager: without file manager interaction, nothing stored in the files can be retrieved

G4

Question: We all use SSOL to choose and register for classes. Another option would be to have a single Google sheet (shared spreadsheet) that we all use to register for classes. What are problems with using a shared spreadsheet?

Answer:

1. For shared spreadsheet, not multiple(too many) users can access and edit it simultaneously. When too many users are editing a shared spreadsheet simultaneously, the effective is decrease and likelihood of errors is increase.
2. Because we are editing a very large shared spreadsheet, a single Google sheet cannot prevent data redundancy.
3. A shared spreadsheet cannot store information as effectively as DBMS
4. A shared spreadshee has record limitation
5. It is not secure for us to use a single Google sheet to register for classes, because anyone who can access this spreadsheet can see personal information of other users of the spreadsheet

G5

Question: NoSQL databases have become increasingly popular for supporting applications. List 3 benefits of or reasons for using NoSQL databases versus SQL/relational databases. List 3 benefits of relational databases versus NoSQL databases.

Answer:

- **using NoSQL databases versus SQL/relational databases:** (1)NoSQL can work with large amounts of unstructured or semi-structured data that doesn't fit the relational model (2)NoSQL allow us to scale-out horizontally, which means we can add cheaper commodity servers whenever we need to (3)NoSQL support flexibility of a dynamic schema or want more choice over the data model
- **using relational databases versus NoSQL databases:** (1)SQL support transaction-oriented systems such as accounting or financial applications (2)SQL enable a high degree of data integrity and security (3)SQL can perform complex queries

Relational Model

R1

Question: A column in a relation (table) has a *type*. Consider implementing a `date` as `CHAR(10)` in the format `YYYY-MM-DD`. The lecture material states that attributes (column values) come from a *domain*. Using `date` explain the difference between a *domain* and a *type*.

Answer:

`Domain` is a set of acceptable values that a column is allowed to contain. This is based on various properties and the data type for the column. `Domain` of `date` is character data in a length of 10 and in the format of YYYY-MM-DD. However data `type` determines what type of values can be stored in a column. `Type` of `date` is character data in a length of 10.

R2

Question: The domain for a relation (table) attribute (column) should be *atomic*. Why?

Answer:

When the domain for a relation attribute is atomic, columns of relational tables aren't repeating groups or arrays. The atomic value property's primary advantage is that it makes data processing logic simpler.

R3

Question: "In the US Postal System, a delivery point is a specific set of digits between 00 and 99 assigned to every address. When combined with the ZIP + 4 code, the delivery point provides a unique identifier for every deliverable address served by the United States Postal Service."

The lecture 2 slides provide a notation for representing a relation's schema. Assume we want to define a relation for US mailing addresses, and that the columns are:

- Zip code
- +4 code
- delivery_point
- address_line_1
- address_line_2
- city
- state

Use the notation to define the schema for an address. A simple example of an address's column values might be:

- Zip code: 10027
- +4 code: 6623
- delivery_point: 99
- address_line_1: 520 W 120th St
- address_line_2: Room 402

- city: New York
- state: NY

Answer:

ADDRESSES(Zip_code, +4_code, delivery_point, address_line_1, address_line_2, city, state)

R4

Note: Use the [RelaX](#) calculator and the schema associated with the recommended textbook to answer this question. Your answer should contain:


- The text for the query.
- An image showing the query execution and result.

An example of the format is:

Query

`σ capacity >= 50 (classroom)`

Execution



| classroom.building | classroom.room_number | classroom.capacity |
|--------------------|-----------------------|--------------------|
| 'Packard' | 101 | 500 |
| 'Taylor' | 3128 | 70 |
| 'Watson' | 120 | 50 |

Question: Translate the following SQL statement into an equivalent relational algebra statement.

```
select
  *
from
  instructor
```

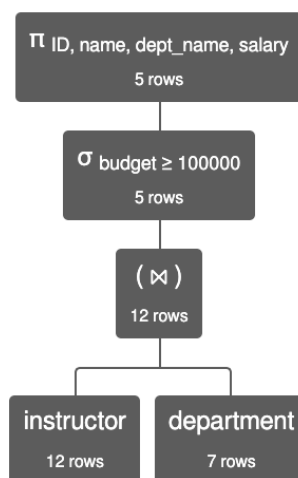
where
dept_name in (select dept_name from department where budget >= 100000)

Answer:

Query

```
 $\pi$  ID, name, dept_name, salary ( $\sigma$  budget >= 100000 (instructor  $\bowtie$  department))
```

Execution



```
 $\pi$  ID, name, dept_name, salary (  $\sigma$  budget  $\geq$  100000 ( instructor  $\bowtie$  department ) )
```

Execution time: 2 ms

| instructor.ID | instructor.name | instructor.dept_name | instructor.salary |
|---------------|-----------------|----------------------|-------------------|
| 10101 | 'Srinivasan' | 'Comp. Sci.' | 65000 |
| 12121 | 'Wu' | 'Finance' | 90000 |
| 45565 | 'Katz' | 'Comp. Sci.' | 75000 |
| 76543 | 'Singh' | 'Finance' | 80000 |
| 83821 | 'Brandt' | 'Comp. Sci.' | 92000 |

R5

Use the same format to answer this question.

Question

Use the following query to compute a new table.

```
section_and_time =  $\pi$  course_id, sec_id, semester, year, day, start_hr, start_min, end_hr,
end_min (section  $\bowtie$  time_slot)
```

Using only section_and_time, write a relational algebra expression that returns a relation of overlapping courses of the form

```
(course_id_1, sec_id_1, semester_1, year_1, course_id_2, sec_id_2, semester_2, year_2) .
```

Your table cannot container duplicates. For example, a result containing

(BIO-101, 1, fall, 2022, MATH-101, 2, fall, 2022)
(MATH-101, 2, fall, 2022, BIO-101, 1, fall, 2022)

is incorrect.

Answer:

Query

```

π A.course_id -> course_id_1, A.sec_id -> sec_id_1, A.semester -> semester_1, A.year ->
year_1,
  B.course_id -> course_id_2, B.sec_id -> sec_id_2, B.semester -> semester_2, B.year ->
year_2
( σ (A.semester = B.semester) ∧ (A.year = B.year) ∧ (A.day = B.day) ∧
  (((A.end_hr = B.start_hr) ∧ (A.end_min > B.start_min)) ∨
   ((A.end_hr > B.start_hr) ∧ (A.end_hr ≤ B.end_hr)) ∨
   ((A.start_hr = B.end_hr) ∧ (A.start_min < B.end_min)) ∨
   ((A.start_hr < B.end_hr) ∧ (A.start_hr ≥ B.start_hr)))
  ((ρA(section_and_time)) ⋈ A.course_id < B.course_id (ρB(section_and_time))))

```

Execution

π A.course_id→course_id_1, A.sec_id→sec_id_1, A.semester→semester_1, A.year→year_1,
B.course_id→course_id_2, B.sec_id→sec_id_2, B.semester→semester_2, B.year→year_2

4 rows

σ (A.semester = B.semester) and (A.year = B.year) and (A.day = B.day) and (((A.end_hr =
B.start_hr) and (A.end_min > B.start_min)) or ((A.end_hr > B.start_hr) and (A.end_hr ≤ B.end_hr))
or ((A.start_hr = B.end_hr) and (A.start_min < B.end_min)) or ((A.start_hr < B.end_hr) and
(A.start_hr ≥ B.start_hr)))

12 rows

(⋈ A.course_id < B.course_id)

768 rows

ρ A

41 rows

ρ B

41 rows

section_and_time π course_id, sec_id, semester, year, day, start_hr, start_min, end_hr,
end_min

41 rows

section_and_time π course_id, sec_id, semester, year, day, start_hr, start_min, end_hr,
end_min

41 rows

(⋈)

41 rows

(⋈)

41 rows

section

15 rows

time_slot

20 rows

section

15 rows

time_slot

20 rows

π A.course_id→course_id_1, A.sec_id→sec_id_1, A.semester→semester_1, A.year→year_1, B.course_id→course_id_2, B.sec_id→sec_id_2, B.semester→semester_2, B.year→year_2 (σ (A.semester = B.semester) and (A.year = B.year) and (A.day = B.day) and (((A.end_hr = B.start_hr) and (A.end_min > B.start_min)) or ((A.end_hr > B.start_hr) and (A.end_hr ≤ B.end_hr)) or ((A.start_hr = B.end_hr) and (A.start_min < B.end_min)) or ((A.start_hr < B.end_hr) and (A.start_hr ≥ B.start_hr))) ((ρ A π course_id, sec_id, semester, year, day, start_hr, start_min, end_hr, end_min (section ⋈ time_slot))) ⋈ A.course_id < B.course_id (ρ B π course_id, sec_id, semester, year, day, start_hr, start_min, end_hr, end_min (section ⋈ time_slot))))

Execution time: 4 ms

| course_id_1 | sec_id_1 | semester_1 | year_1 | course_id_2 | sec_id_2 | semester_2 | year_2 |
|-------------|----------|------------|--------|-------------|----------|------------|--------|
| 'CS-319' | 1 | 'Spring' | 2010 | 'FIN-201' | 1 | 'Spring' | 2010 |
| 'CS-319' | 2 | 'Spring' | 2010 | 'HIS-351' | 1 | 'Spring' | 2010 |
| 'CS-315' | 1 | 'Spring' | 2010 | 'MU-199' | 1 | 'Spring' | 2010 |
| 'CS-347' | 1 | 'Fall' | 2009 | 'PHY-101' | 1 | 'Fall' | 2009 |

<

1

>

SQL

- You will use the [Classic Models tutorial database](#), which you should have already loaded into MySQL.

S1

Question: Create a view `employee_customer_sales` with the following information:

- `employeeNumber`

localhost:8888/nbconvert/html/Desktop/f22_w4111_final_exam/F22_W4111_003_Final_Exam.ipynb?download=false

7/17

- `employeeLastname`
- `employeeFirstName`
- `customerNumber`
- `customerName`
- `revenue`
- The employee information is for the employee that is the `customer.customerRepEmployeeNumber`.
- `revenue` is the total revenue over all of the customer's orders.
 - The revenue for an `order` is `priceEach*quantityOrdered` for each `orderdetails` in the order.

Answer:

In [12]:

```
%%sql

USE classicmodels;

CREATE OR REPLACE VIEW employee_customer_sales AS
WITH order_revenue AS (
  SELECT customerNumber,
         orders.orderNumber,
         SUM(quantityOrdered * priceEach) AS order_revenue
  FROM orderdetails
  INNER JOIN orders
    ON orderdetails.orderNumber = orders.orderNumber
  GROUP BY orderNumber
)
SELECT employeeNumber AS employeeNumber,
       lastName AS employeeLastname,
       firstName AS employeeFirstName,
       order_revenue.customerNumber AS customerNumber,
       customerName AS customerName,
       SUM(order_revenue) AS revenue
FROM order_revenue
INNER JOIN customers
  ON order_revenue.customerNumber = customers.customerNumber
INNER JOIN employees
  ON employees.employeeNumber = customers.salesRepEmployeeNumber
GROUP BY customerNumber
ORDER BY employeeLastname, revenue DESC;
```

```
* mysql+pymysql://root:***@localhost
0 rows affected.
0 rows affected.
```

Out[12]: []

Test Answer:

In [13]:

```
%sql select * from employee_customer_sales;
```

```
* mysql+pymysql://root:***@localhost
98 rows affected.
```

Out[13]:

| employeeNumber | employeeLastname | employeeFirstName | customerNumber | customerName | revenue |
|----------------|------------------|-------------------|----------------|---------------------------|-----------|
| 1337 | Bondur | Loui | 146 | Saveley & Henriot, Co. | 130305.35 |
| 1337 | Bondur | Loui | 353 | Reims Collectables | 126983.19 |
| 1337 | Bondur | Loui | 172 | La Corne D'abondance, Co. | 86553.52 |
| 1337 | Bondur | Loui | 406 | Auto Canal+ Petit | 86436.97 |
| 1337 | Bondur | Loui | 350 | Marseille Mini Autos | 71547.53 |
| 1337 | Bondur | Loui | 250 | Lyon Souvenirs | 67659.19 |
| 1501 | Bott | Larry | 187 | AV Stores, Co. | 148410.09 |
| 1501 | Bott | Larry | 201 | UK Collectables, Ltd. | 106610.72 |
| 1501 | Bott | Larry | 334 | Suominen Souvenirs | 103896.74 |
| 1501 | Bott | Larry | 311 | Oulu Toy Supplies, Inc. | 95706.15 |

| | | | | | |
|------|-----------|--------|-----|--------------------------------|-----------|
| 1501 | Bott | Larry | 186 | Toys of Finland, Co. | 95546.46 |
| 1501 | Bott | Larry | 324 | Stylish Desk Decors, Co. | 80556.73 |
| 1501 | Bott | Larry | 240 | giftsbymail.co.uk | 71783.75 |
| 1501 | Bott | Larry | 489 | Double Decker Gift Stores, Ltd | 29586.15 |
| 1401 | Castillo | Pamela | 382 | Salzburg Collectables | 137480.07 |
| 1401 | Castillo | Pamela | 145 | Danish Wholesale Imports | 129085.12 |
| 1401 | Castillo | Pamela | 278 | Rovelli Gifts | 127529.69 |
| 1401 | Castillo | Pamela | 386 | L'ordine Souvenirs | 125505.57 |
| 1401 | Castillo | Pamela | 227 | Heintze Collectables | 89909.80 |
| 1401 | Castillo | Pamela | 249 | Amica Models & Co. | 82223.23 |
| 1401 | Castillo | Pamela | 314 | Petit Auto | 70851.58 |
| 1401 | Castillo | Pamela | 452 | Mini Auto Werke | 51059.99 |
| 1401 | Castillo | Pamela | 381 | Royale Belge | 29217.18 |
| 1401 | Castillo | Pamela | 473 | Frau da Collezione | 25358.32 |
| 1188 | Firrelli | Julie | 320 | Mini Creations Ltd. | 101872.52 |
| 1188 | Firrelli | Julie | 379 | Collectables For Less Inc. | 73533.65 |
| 1188 | Firrelli | Julie | 495 | Diecast Collectables | 65541.74 |
| 1188 | Firrelli | Julie | 339 | Classic Gift Ideas, Inc | 57939.34 |
| 1188 | Firrelli | Julie | 204 | Online Mini Collectables | 55577.26 |
| 1188 | Firrelli | Julie | 173 | Cambridge Collectables Co. | 32198.69 |
| 1611 | Fixter | Andy | 114 | Australian Collectors, Co. | 180585.07 |
| 1611 | Fixter | Andy | 276 | Anna's Decorations, Ltd | 137034.22 |
| 1611 | Fixter | Andy | 282 | Souvenirs And Things Co. | 133907.12 |
| 1611 | Fixter | Andy | 471 | Australian Collectables, Ltd | 55866.02 |
| 1611 | Fixter | Andy | 333 | Australian Gift Network, Co | 55190.16 |
| 1702 | Gerard | Martin | 458 | Corrida Auto Replicas, Ltd | 112440.09 |
| 1702 | Gerard | Martin | 298 | Vida Sport, Ltd | 108777.92 |
| 1702 | Gerard | Martin | 216 | Enaco Distributors | 68520.47 |
| 1702 | Gerard | Martin | 484 | Iberia Gift Imports, Corp. | 50987.85 |
| 1702 | Gerard | Martin | 344 | CAF Imports | 46751.14 |
| 1370 | Hernandez | Gerard | 141 | Euro+ Shopping Channel | 820689.54 |
| 1370 | Hernandez | Gerard | 119 | La Rochelle Gifts | 158573.12 |
| 1370 | Hernandez | Gerard | 209 | Mini Caravy | 75859.32 |
| 1370 | Hernandez | Gerard | 171 | Daedalus Designs Imports | 61781.70 |
| 1370 | Hernandez | Gerard | 242 | Alpha Cognac | 60483.36 |
| 1370 | Hernandez | Gerard | 256 | Auto Associés & Cie. | 58876.41 |
| 1370 | Hernandez | Gerard | 103 | Atelier graphique | 22314.36 |
| 1165 | Jennings | Leslie | 124 | Mini Gifts Distributors Ltd. | 591827.34 |
| 1165 | Jennings | Leslie | 450 | The Sharp Gifts Warehouse | 143536.27 |
| 1165 | Jennings | Leslie | 321 | Corporate Gift Ideas Co. | 132340.78 |
| 1165 | Jennings | Leslie | 161 | Technics Stores Inc. | 104545.22 |
| 1165 | Jennings | Leslie | 129 | Mini Wheels Co. | 66710.56 |
| 1165 | Jennings | Leslie | 487 | Signal Collectibles Ltd. | 42570.37 |
| 1504 | Jones | Barry | 448 | Scandinavian Gift Ideas | 120943.53 |
| 1504 | Jones | Barry | 121 | Baane Mini Imports | 104224.79 |

| | | | | | |
|------|-----------|----------|-----|------------------------------------|-----------|
| 1504 | Jones | Barry | 167 | Herkku Gifts | 97562.47 |
| 1504 | Jones | Barry | 259 | Toms Spezialitäten, Ltd | 89223.14 |
| 1504 | Jones | Barry | 128 | Blauer See Auto, Co. | 75937.76 |
| 1504 | Jones | Barry | 299 | Norway Gifts By Mail, Co. | 69059.04 |
| 1504 | Jones | Barry | 144 | Volvo Model Replicas, Co | 66694.82 |
| 1504 | Jones | Barry | 189 | Clover Collections, Co. | 49898.27 |
| 1504 | Jones | Barry | 415 | Bavarian Collectables Imports, Co. | 31310.09 |
| 1612 | Marsh | Peter | 323 | Down Under Souvenirs, Inc | 154622.08 |
| 1612 | Marsh | Peter | 496 | Kelly's Gift Shop | 137460.79 |
| 1612 | Marsh | Peter | 166 | Handji Gifts& Co | 107746.75 |
| 1612 | Marsh | Peter | 357 | GiftsForHim.com | 94431.76 |
| 1612 | Marsh | Peter | 412 | Extreme Desk Decorations, Ltd | 90332.38 |
| 1621 | Nishi | Mami | 148 | Dragon Souvenirs, Ltd. | 156251.03 |
| 1621 | Nishi | Mami | 398 | Tokyo Collectables, Ltd | 105548.73 |
| 1621 | Nishi | Mami | 385 | Cruz & Sons Co. | 87468.30 |
| 1621 | Nishi | Mami | 177 | Osaka Souvenirs Co. | 62361.22 |
| 1621 | Nishi | Mami | 211 | King Kong Collectables, Co. | 45480.79 |
| 1216 | Patterson | Steve | 363 | Online Diecast Creations Co. | 116449.29 |
| 1216 | Patterson | Steve | 157 | Diecast Classics Inc. | 104358.69 |
| 1216 | Patterson | Steve | 286 | Marta's Replicas Co. | 90545.37 |
| 1216 | Patterson | Steve | 462 | FunGiftIdeas.com | 88627.49 |
| 1216 | Patterson | Steve | 362 | Gifts4AllAges.com | 84340.32 |
| 1216 | Patterson | Steve | 198 | Auto-Moto Classics Inc. | 21554.26 |
| 1166 | Thompson | Leslie | 205 | Toys4GrownUps.com | 93803.30 |
| 1166 | Thompson | Leslie | 239 | Collectable Mini Designs Co. | 80375.24 |
| 1166 | Thompson | Leslie | 112 | Signal Gift Stores | 80180.98 |
| 1166 | Thompson | Leslie | 475 | West Coast Collectables Co. | 43748.72 |
| 1166 | Thompson | Leslie | 347 | Men 'R' US Retailers, Ltd. | 41506.19 |
| 1166 | Thompson | Leslie | 219 | Boards & Toys Co. | 7918.60 |
| 1286 | Tseng | Foon Yue | 151 | Muscle Machine Inc | 177913.95 |
| 1286 | Tseng | Foon Yue | 181 | Vitachrome Inc. | 72497.64 |
| 1286 | Tseng | Foon Yue | 455 | Super Scale Inc. | 70378.65 |
| 1286 | Tseng | Foon Yue | 424 | Classic Legends Inc. | 69214.33 |
| 1286 | Tseng | Foon Yue | 233 | Québec Home Shopping Network | 68977.67 |
| 1286 | Tseng | Foon Yue | 456 | Microscale Inc. | 29230.43 |
| 1323 | Vanauf | George | 131 | Land of Toys Inc. | 149085.15 |
| 1323 | Vanauf | George | 175 | Gift Depot Inc. | 95424.63 |
| 1323 | Vanauf | George | 328 | Tekni Collectables Inc. | 81806.55 |
| 1323 | Vanauf | George | 319 | Mini Classics | 78432.16 |
| 1323 | Vanauf | George | 486 | Motor Mint Distributors Inc. | 77726.59 |
| 1323 | Vanauf | George | 202 | Canadian Gift Exchange Network | 70122.19 |
| 1323 | Vanauf | George | 260 | Royal Canadian Collectables, Ltd. | 66812.00 |
| 1323 | Vanauf | George | 447 | Gift Ideas Corp. | 49967.78 |

Question:

- Below, there is a query that creates a view. Run the query.
- Using the view, write a query that produces a table of the form `(productCode, productName)` for products that no customer in Asia has ordered.
- For this questions purposes, the Asian countries are:
 - Japan
 - Singapore
 - Philipines
 - Hong King
- You must not use a JOIN.

In [14]:

```
#
# Create the view
#
%sql create or replace view orders_all as \
    select * from orders join orderdetails using(orderNumber)

* mysql+pymysql://root:***@localhost
0 rows affected.
```

Out[14]: []

Answer:

In [15]:

```
%%sql

WITH asian_code AS (
    SELECT productCode
    FROM orders_all
    WHERE customerNumber IN (
        SELECT customerNumber
        FROM customers
        WHERE country = 'Japan' OR
            country = 'Singapore' OR
            country = 'Philippines' OR
            country = 'Hong Kong'
    )
)

SELECT productCode,
       productName
FROM products
WHERE productCode NOT IN (SELECT productCode FROM asian_code)

* mysql+pymysql://root:***@localhost
15 rows affected.
```

Out[15]:

| productCode | productName |
|-------------|---|
| S10_1678 | 1969 Harley Davidson Ultimate Chopper |
| S10_4757 | 1972 Alfa Romeo GTA |
| S12_2823 | 2002 Suzuki XREO |
| S18_1342 | 1937 Lincoln Berline |
| S18_1367 | 1936 Mercedes-Benz 500K Special Roadster |
| S18_2795 | 1928 Mercedes-Benz SSK |
| S18_2870 | 1999 Indy 500 Monte Carlo SS |
| S18_3029 | 1999 Yamaha Speed Boat |
| S18_3233 | 1985 Toyota Supra |
| S18_3320 | 1917 Maxwell Touring Car |
| S18_3856 | 1941 Chevrolet Special Deluxe Cabriolet |
| S24_2022 | 1938 Cadillac V-16 Presidential Limousine |
| S24_2972 | 1982 Lamborghini Diablo |

S24_4258

1936 Chrysler Airflow

S700_3505

The Titanic

S3

Question:

- Use the `customers` and `orders` for this query.
- Shipping days is the number of days between `orderDate` and `shippedDate`.
- Product a table of the form:
 - `customerNumber`
 - `customerName`
 - `noOfOrders` is the number of orders the customer placed.
 - `averageShippingDays`, which is the average shipping days.
 - `minimumShippingDays`, which is the minimum shipping days.
 - `maximumShippingDays`, which is the maximum shipping days.
- The table should only contain entries where:
 - `noOfOrders >= 3`
 - `averageShippingDays >= 5` or `maximumShippingDays >= 10`.

Answer:

In [16]:

```
%%sql

SELECT customers.customerNumber,
       customerName,
       COUNT(*) AS noOfOrders,
       AVG(datediff(shippedDate, orderDate)) AS averageShippingDays,
       MIN(datediff(shippedDate, orderDate)) AS minimumShippingDays,
       MAX(datediff(shippedDate, orderDate)) AS maximumShippingDays
FROM customers
INNER JOIN orders
  ON customers.customerNumber = orders.customerNumber
GROUP BY customers.customerNumber
HAVING noOfOrders >= 3 AND
       (averageShippingDays >= 5 OR maximumShippingDays >= 10)
```

```
* mysql+pymysql://root:***@localhost
12 rows affected.
```

Out [16]:

| | customerNumber | customerName | noOfOrders | averageShippingDays | minimumShippingDays | maximumShippingDays |
|--|----------------|------------------------------|------------|---------------------|---------------------|---------------------|
| | 363 | Online Diecast Creations Co. | 3 | 5.0000 | 4 | 6 |
| | 385 | Cruz & Sons Co. | 3 | 5.3333 | 5 | 6 |
| | 148 | Dragon Souvenirs, Ltd. | 5 | 14.6000 | 1 | 65 |
| | 198 | Auto-Moto Classics Inc. | 3 | 5.6667 | 5 | 6 |
| | 161 | Technics Stores Inc. | 4 | 5.2500 | 4 | 6 |
| | 205 | Toys4GrownUps.com | 3 | 5.3333 | 4 | 6 |
| | 276 | Anna's Decorations, Ltd | 4 | 5.0000 | 4 | 6 |
| | 462 | FunGiftIdeas.com | 3 | 5.0000 | 3 | 6 |
| | 448 | Scandinavian Gift Ideas | 3 | 5.5000 | 5 | 6 |
| | 328 | Tekni Collectables Inc. | 3 | 5.0000 | 4 | 6 |
| | 209 | Mini Caravy | 3 | 5.6667 | 5 | 6 |
| | 398 | Tokyo Collectables, Ltd | 4 | 5.5000 | 2 | 8 |

Graph Database — Neo4j

- You will use your online/cloud Neo4j database for these problems.
- You must have loaded the Movie sample data.

N1

Question:

- The relationship `REVIEWED` connects a `Person` and `Movie`, and has the properties `rating` and `summary`.
- Write Python code using `py2neo` that produces the following table.

In [17]: `import pandas as pd`

Answer:

In [18]:

```
graph = Graph(neo4j_url, auth=(neo4j_user, neo4j_password))
q = "MATCH p=(n:Person)-[r:REVIEWED]->(m:Movie) return n.name , r.rating, r.summary, m.title"
result = graph.run(q)
df = pd.DataFrame(result,columns=['reviewer_name', 'rating', 'rating_summary', 'movie_title'])
df.sort_values(by=['reviewer_name'],ignore_index=True)
```

Out[18]:

| | reviewer_name | rating | rating_summary | movie_title |
|---|------------------|--------|---|-------------------|
| 0 | Angela Scope | 62 | Pretty funny at times | The Replacements |
| 1 | James Thompson | 100 | The coolest football movie ever | The Replacements |
| 2 | James Thompson | 65 | Fun, but a little far fetched | The Da Vinci Code |
| 3 | Jessica Thompson | 92 | You had me at Jerry | Jerry Maguire |
| 4 | Jessica Thompson | 65 | Silly, but fun | The Replacements |
| 5 | Jessica Thompson | 45 | Slapstick redeemed only by the Robin Williams ... | The Birdcage |
| 6 | Jessica Thompson | 85 | Dark, but compelling | Unforgiven |
| 7 | Jessica Thompson | 95 | An amazing journey | Cloud Atlas |
| 8 | Jessica Thompson | 68 | A solid romp | The Da Vinci Code |

N2

Question:

- There are relationships `ACTED_IN` and `DIRECTED` between `Person` and `Movie`.
- Write Python code that produces the following table that shows people or both acted in and directed a movie.

In [19]:

```
graph = Graph(neo4j_url, auth=(neo4j_user, neo4j_password))
q = "MATCH p=(a:Person)-[:ACTED_IN]->(m:Movie)<-[:DIRECTED]-(b:Person) WHERE a=b return a.name, m.title"
result = graph.run(q)
df = pd.DataFrame(result,columns=['Name', 'Title'])
df
```

Out[19]:

| | Name | Title |
|---|----------------|-------------------|
| 0 | Tom Hanks | That Thing You Do |
| 1 | Clint Eastwood | Unforgiven |
| 2 | Danny DeVito | Hoffa |

MongoDB

- Run the following code using your Atlas MongoDB.

```
In [20]: import json

client = pymongo.MongoClient(
    mongodb_url
)

with open("episodes.json") as e_file:
    episodes = json.load(e_file)

episodes = episodes['episodes']

for e in episodes:
    client['w4111_final']['episodes'].insert_one(e)
```

```
In [21]: ratings_df = pd.read_csv("got_title_ratings.csv")
ratings_info = ratings_df[['tconst', 'averageRating', 'numVotes']]
r_dict = ratings_info.to_dict("records")

for r in r_dict:
    client['w4111_final']['ratings'].insert_one(r)
```

Question:

Write Python code that uses an aggregation pipeline and operations to produce the following table.

```
In [22]: # Requires the PyMongo package.
# https://api.mongodb.com/python/current
#
# Write the query/aggregation that produces result
```

```
In [23]: collection_episodes = client['w4111_final']['episodes']

result = collection_episodes.aggregate( [
    {
        '$set': {
            'episodeLink': { '$substr': ['$episodeLink', 7, 9] }
        }
    },
    {
        '$lookup':
        {
            'from': 'ratings',
            'localField': 'episodeLink',
            'foreignField': 'tconst',
            'as': 'rating_info'
        }
    },
    {
        '$addFields': {
            'avgRating': { '$arrayElemAt': ['$rating_info.averageRating', 0] },
            'numVotes': { '$arrayElemAt': ['$rating_info.numVotes', 0] }
        }
    },
    {
        '$project': {
            'seasonNum': 1,
            'episodeNum': 1,
            'episodeLink': 1,
            'episodeTitle': 1,
            'avgRating': 1,
            'numVotes': 1
        }
    }
])

info_df = pd.DataFrame(list(result))
info_df = info_df[['seasonNum', 'episodeNum', 'episodeLink', 'episodeTitle', 'avgRating', 'numVotes']]
info_df
```

Out [23]:

| | seasonNum | episodeNum | episodeLink | episodeTitle | avgRating | numVotes |
|-----|-----------|------------|-------------|---------------------------------------|-----------|----------|
| 0 | 1 | 1 | tt1480055 | Winter Is Coming | 8.9 | 48686 |
| 1 | 1 | 2 | tt1668746 | The Kingsroad | 8.6 | 36837 |
| 2 | 1 | 3 | tt1829962 | Lord Snow | 8.5 | 34863 |
| 3 | 1 | 4 | tt1829963 | Cripples, Bastards, and Broken Things | 8.6 | 33136 |
| 4 | 1 | 5 | tt1829964 | The Wolf and the Lion | 9.0 | 34436 |
| ... | ... | ... | ... | ... | ... | ... |
| 68 | 8 | 2 | tt6027908 | A Knight of the Seven Kingdoms | 7.9 | 130844 |
| 69 | 8 | 3 | tt6027912 | The Long Night | 7.5 | 215995 |
| 70 | 8 | 4 | tt6027914 | The Last of the Starks | 5.5 | 165067 |
| 71 | 8 | 5 | tt6027916 | The Bells | 6.0 | 192449 |
| 72 | 8 | 6 | tt6027920 | The Iron Throne | 4.0 | 248318 |

73 rows x 6 columns

Data Modeling and Schema Definition

- This is an exciting, interesting problem that involves:
 - Using Crow's Foot Notation
 - Relational approaches to implementing specialization, aggregation, quaternary relations, composite attributes and multi-valued attributes.
 - Foreign keys, check constraints and triggers.
- I did the answer and it took 3 hours to do all the work. My normal rule of thumb is that students require about 15 times as much time as I need to produce an answer.
- I giggled like the Riddler in Batman about how much fun we were going to have working on this question, and then the following happened.

DO NOT ASSIGN THAT QUESTION!
THE EXAM IS INSANELY LONG ALREADY!
AND THEY HAVE ALREADY DONE DATA MODELING!



- So, there will not be any data modeling question on the exam. Darn!

Module II Questions

- The questions require brief, written answers.

Q1

Question:

Briefly explain:

- Functional Dependency
- Lossy Decomposition
- Normalization

Answer:

1. Functional Dependencies: It is a constraint in a DBMS that establishes the relationship between one attribute and another attribute.
2. Lossy Decomposition: It is the process of a relation be divided into several relational schemas, which results in information being lost when the original relation is retrieved.
3. Normalization: In order to reduce redundancy from a relation or group of relations, normalization splits the larger table into smaller ones and connects them using relationships.

Q2

Question:

Briefly explain:

- Serializability
- Conflict Serializability
- Deadlock
- Cascading Abort
- Two Phase Locking

Answer:

1. Serializability: The primary accuracy test for concurrent transaction executions is serializability. It is regarded as the highest level of transaction isolation and is crucial to concurrency control.
2. Conflict Serializability: a schedule is conflict serializable if it is conflict equivalent to a serial schedule
3. Deadlock: It is a case that two or more transactions are waiting indefinitely for one another to give up locks
4. Cascading Abort: In order to stop the second transaction from accessing invalid data, one transaction must terminate in order for the other to continue.
5. Two Phase Locking: Whenever Locking and Unlocking may be completed in two phases: 1) Growing Phase: Although no locks on data items can be released, new ones may be obtained. 2) Shrinking Phase: Only existing locks may be released during this time; no new locks may be purchased.

Q3

Question:

Briefly explain:

- Logical block addressing, CHS addressing
- RAID-0, RAID-1, RAID-5

- Fixed length records, variable length records.

Answer:

1. Logical block addressing: It is a method for identifying the position of data blocks stored on computer storage devices, typically secondary storage systems like hard disk drives.
2. CHS addressing: It is the method of locating specific disk sectors based on where they are located within a track, where the track is identified by the head and cylinder numbers.
3. RAID-0: it requires at least two drives, we can combine the drives and write data to both of them at once or sequentially, depending on the system.
4. RAID-1: RAID-1 will duplicate data and save a copy on each drive if we have at least two drives. Mirroring prevents file loss in the event of a drive failure.
5. RAID-5: Using RAID-5 will divide data into segments and save those parts across drives if we have at least three hard drives.
6. Fixed length records: the length of the fields in each record has been set to be a certain maximum number of characters long.
7. Variable length records: the length of a field can change to allow data of any size to fit.

Q4

Question:

Briefly explain:

- Clustered Index
- Sparse Index
- Covering Index

Answer:

1. Clustered Index: the index whose search key specifies the sequential order of the file.
2. Sparse Index: contains index records for only some search-key values.
3. Covering Index: an index covers the query if all the columns specified in the query are part of the index

Q5

Question:

Briefly explain:

- Equivalent queries
- Hash Join
- Materialization, Pipelining

Answer:

1. Equivalent queries: queries that will produce the exactly the same answer.
2. Hash join: It is a way of executing a join where a hash table is used to find matching rows between one or more tables
3. Materialization: generate results of an expression whose inputs are relations or are already computed, store it on disk, then repeat
4. Pipelining: pass on tuples to parent operations even the operation is still in progress